

LMMT-SecureDove

CPTS 428 Fall 2023
Professor Haipeng Cai

Logan Kloft
Manjesh Puram
Matthew Gerola
Taylor West



Table of Contents

Table of Contents.....	2
Introduction.....	3
Confidentiality.....	4
Integrity.....	10
Availability.....	13
Improved Security Plan.....	20

Introduction

In the remainder of the report, you will see security tests targeting the confidentiality, integrity, and availability of our application. These tests target our original security goals, in addition to other security vulnerabilities that we have thought of since the security goal / measurement report. We signify failed tests using  which means there exists a vulnerability. We indicate successful tests using  which means our project implements security measurements that guard against the targeted vulnerability.

After these tests, we summarize the findings of the break / attack part and enhance the security plan to reflect any vulnerabilities discovered during the test. You may find code used to test security requirements in our repository linked here:

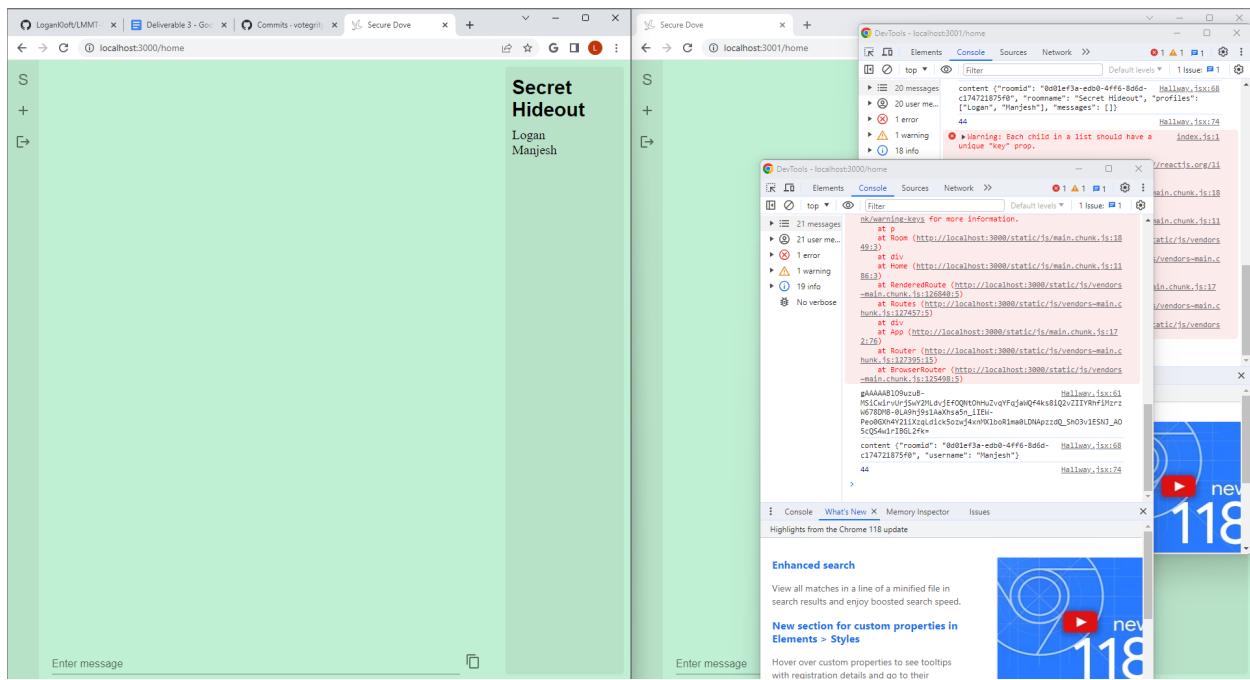
<https://github.com/LoganKloft/LMMT-SecureDove>. Also, for each test that has code associated with it, we link to the code for convenience.

Confidentiality

Goal: At every step of message transferral, the message is encrypted 

How: Trace message path from creation in a client to receipt in another client

We will trace the route of one client sending a message to a chat room and then checking that it is encrypted when received by the server and also encrypted when broadcasted from the server (received on the client). To do this, we have embedded print statements in both the backend and frontend. Both the backend and frontend only print the “content” of the messages they receive. The content carries the actual payload of information that we wish to keep confidential. We use the developer tools to monitor print statements on the client:



We use a terminal to view the backend's output:

```
who: Manjesh
{'type': 'room', 'verb': 'get', 'id': '84721251-b651-4def-bed1-7e3b9eb4c80', 'content': {'roomid': '0d01ef3a-edb0-4ff6-8d6d-c174721875f0'}}

===== SINGLE RESPONSE =====
{'type': 'room', 'verb': 'get', 'content': {'roomid': '0d01ef3a-edb0-4ff6-8d6d-c174721875f0', 'roomname': 'Secret Hideout', 'profiles': ['Logan', 'Manjesh'], 'messages': []} }

IN PARSER
gAAAAAB10u1nINIF2-srqFeMTveGwmI-YrxJJ44dnyMUpNn90PvRdBnFoEv_MsyC_NDhX4Lfx8bdj1B44K7kHcqgXVF2__wFC6nlhtj-YUvrfXBKteK6RwXNGUFeTnP1tCf9_Go7
FINISH DECODE
{'type': 'ACK', 'verb': 'post', 'content': {'id': '9c485ae1-cdc1-4717-b82b-aa9d03b2a5f7'}}}
```

We start by entering: “This is a private message” on the left client (Logan). We will observe what the server receives and verify that the “content” is encrypted.

```

IN PARSER
gAAAAAB1O9zFXTANTrVVrPkRH-CtP7KHSmCOLiwFOf0RTrgv_0C2qBET1MwiTDHb5P9KrXLwO6-0g4V5A1gWxxSfD1lV3e2kiyKLGfs6F8prZeTCwZhRqk3Jcu01MnuvAjh3kd1ksu0yykfPT
ml05aoPjmFosEg1gFss-o87UgoB5EVZHswws5b_TDt8ZHYJ1NC_yYbmBw
FINISH DECODE
{ 'type': 'message', 'verb': 'post', 'id': '07cc8101-c292-430a-a8c0-832553adb7d8', 'content': { 'roomid': '0d01ef3a-edb0-4ff6-8d6d-c174721875f0', 'message': 'This is a private message'} }
===== SEND ACK ENCRYPTED =====
{ 'type': 'ACK', 'verb': 'post', 'content': { 'id': '07cc8101-c292-430a-a8c0-832553adb7d8'} }

===== REQUEST =====
who: Logan
{ 'type': 'message', 'verb': 'post', 'id': '07cc8101-c292-430a-a8c0-832553adb7d8', 'content': { 'roomid': '0d01ef3a-edb0-4ff6-8d6d-c174721875f0', 'message': 'This is a private message'} }

```

In the above image we see that the content looks like a bunch of random characters (the stuff under “IN PARSER”). But once decrypting the content, then printing the entire request, we see that the message was successfully encrypted and transported to the server (the “FINISH DECODE” and “== REQUEST ==”). Now we need to confirm that the message is still encrypted when received by clients in the chatroom.

The screenshot shows two browser developer tool Network tabs side-by-side. Both tabs are for the URL <http://localhost:3001/static/js/vendors~main.chunk.js:125498:5>. The left tab is labeled "at BrowserRouter". The right tab is also labeled "at BrowserRouter". Both tabs show a single request labeled "44". The response body in both tabs is identical:

```

at BrowserRouter (http://localhost:3001/static/js/vendors~main.chunk.js:125498:5)
44
{
  "roomid": "0d01ef3a-edb0-4ff6-8d6d-c174721875f0",
  "header": "Logan @ 27/10/2023 08:52:36",
  "message": "This is a private message"
}

```

Here we can see that both clients receive encrypted content. Then upon decrypting the content and printing the entire response object, we see that it is indeed the original message that was sent from ‘Logan’. This demonstration shows that the message is encrypted during transport, so an attacker would have to either gain access to the server or a client to view unencrypted message content.

Goal: The message is undecipherable while encrypted ✘

How: Analyze complexity of the algorithm

We use symmetric encryption on two separate branches. On the main branch we use fernet symmetric encryption. On the asymmetric branch, we use aes-256-ctr. This algorithm is equivalent to AES with a 256-bit key except that it also includes an “iv” vector for additional security. This makes it safe against CPAs (chosen-plaintext attacks). Work in the future might include picking the iv vector more carefully; currently it’s chosen with random bytes and it’s important to avoid using the same iv for multiple messages. It could be that the crypto library we use has a cryptographically secure randomBytes function, but we would have to do more research.

For the main branch, after a user is created, all “content” is encrypted. However, there is an inherent vulnerability of someone stealing the symmetric key during the initial client-server transaction that shares the symmetric key. The following is a screenshot of the symmetric key with public visibility:

```
[{"type": "profile", "verb": "post", "id": "a20ff3e5-d51e-49bc-8198-73327c8c19e8", "content": {"username": "Logan", "public": "-----BEGIN PUBLIC KEY-----\nMIICjANBgkqhkiG9w0BAQEFAOCg8AMIICCgKCAgEAjN6/wwMMu0pVe9k1OvQ)XZPEbKujhB5e...\n-----END PUBLIC KEY-----"}, "type": "profile", "verb": "post", "content": {"username": "Logan", "userid": "723d9c7e-e554-4795-acf7-7e83ad1e4e51", "symmetric": "QQLUVd9g1pwquj2nwC8jM7QyZqMpXHUNNvzMyOubm1I="}, "id": "a63a53a6-2990-4739-9087-136f1c979ff1"}]
```

The security improvements suggest a solution to prevent the symmetric key from being public observable. Assuming we are able to secure the symmetric key, let's analyze the complexity of the algorithm to decide whether Fernet is susceptible to brute-force attacks. Fernet uses AES in CBC mode with a 128-bit key for encryption with PKCS7 padding¹. Estimates for brute forcing AES-128-CBC are anywhere from ~3400 years to 2.158×10^9 years². Clearly, brute forcing AES is not feasible.

1. <https://cryptography.io/en/latest/fernet/#limitations>
2. [Estimating aes brute force](#)

Goal: Chat rooms are not able to be easily joined ✓

How: Calculate probability of guessing a chatroom id

Chat Room ids are used to join a chat room. They are generated using python's uuid library. Version 4 UUIDs are 128 bits long and reserve 2 bits for the UUID variant and 4 for the UUID version¹. This leaves 122 bits to be randomly generated. Thus there are 2^{122} or 5.3 undecillion (5.3×10^{36}) unique version 4 UUIDs².

Let's see how long it would take to have a 50% chance of guessing the uuid. One would need to generate 2^{121} unique uuids (half of the number of unique version 4 UUIDs). If we assume that a user may generate 1 billion uuids per second, then it would take them $2^{121} / 1 \text{ billion} = 2.658456 \times 10^{27}$ seconds. In 365 day years that is $2.658456 \times 10^{27} / 31,536,000 = 8.4299086 \times 10^{19}$ years. Not to mention the time of checking whether the uuid is correct will take longer than actually generating the uuid. It becomes clear that it is unreasonable to guess a room's uuid.

As an extension, we tested generating uuids on the following processor: Intel(R) Core(TM) i7-10510U CPU @ 1.80GHz 2.30 GHz. Here is an example output from the program:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS GITLENS

PS C:\Users\lklof\OneDrive\Documents\Fall 2023 Classes\CPTS 428\LMMT-SecureDove\Documents\Milestone 3> python .\uuidv4test.py
The execution time for 1,000,000 uuid is: 1.588454246520996
PS C:\Users\lklof\OneDrive\Documents\Fall 2023 Classes\CPTS 428\LMMT-SecureDove\Documents\Milestone 3>
```

The program linked is here:

<https://github.com/LoganKloft/LMMT-SecureDove/blob/main/Documents/Milestone%203/uuidv4test.py>

We can see that it took about 1.588 seconds to generate 1 million version 4 UUIDs. Now multiply the time by 1000 to estimate how long it would take to generate 1 billion UUIDs. It would take about 26.467 minutes. Clearly the idea of generating 1 billion UUIDs per second is

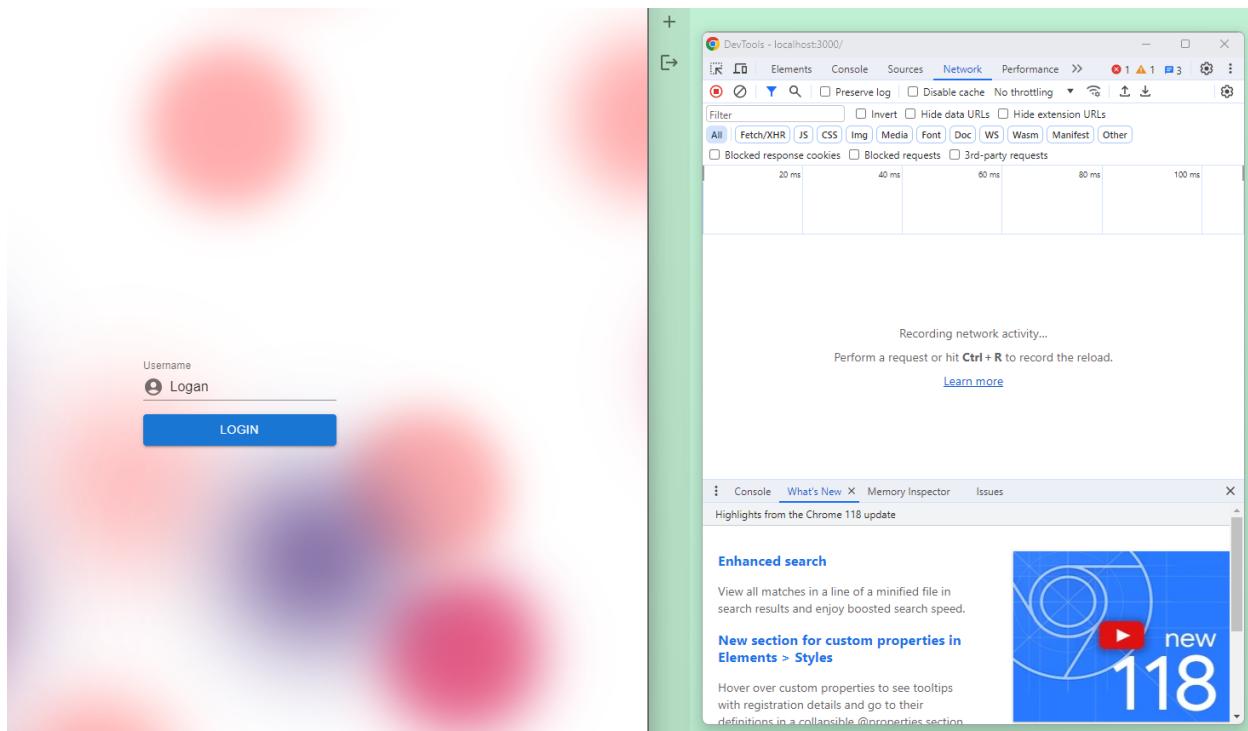
not feasible for a common computer, so to attempt the feat of guessing a UUID, one would need vast compute resources.

1. [RFC 4122](#)
2. [Universally Unique Identifier](#)

How: Intercept a chatroom creation / join request using developer tools

Now, instead of trying to guess a uuid, how about intercepting it? In the following we use the network developer tools to observe outbound requests by one user to see whether the room id is publicly visible.

This image shows the network tab



After logging in, we see the websocket connection:

The screenshot shows the Network tab in the Chrome DevTools interface. At the top, there are several filter options: 'Preserve log' (unchecked), 'Disable cache' (unchecked), 'No throttling' (selected), and a dropdown menu set to 'Normal'. Below these are buttons for 'Invert' (unchecked), 'Hide data URLs' (unchecked), and 'Hide extension URLs' (unchecked). The main table has columns for Name, Status, Type, Initiator, Size, Time, and Waterfall. A single row is visible, representing a request to 'localhost' with status 101, type 'webs...', initiator 'Login.jsx:22', size 0 B, and time pending. At the bottom, it shows '1 requests | 0 B transferred | 0 B resources'.

Name	Status	Type	Initiator	Size	Time	Waterfall
localhost	101	webs...	Login.jsx:22	0 B	Pendi...	

1 requests | 0 B transferred | 0 B resources

Console What's New Memory Inspector Issues

Enhanced search

View all matches in a line of a minified file in search results and enjoy boosted search speed.

New section for custom properties in Elements > Styles

Hover over custom properties to see tooltips with registration details and go to their [definitions in a collapsible @properties section](#).



We click the request and view the messages pane:

Screenshot of the Network tab in the Chrome DevTools Network panel. The 'Messages' tab is selected. A single request to 'localhost' is shown, with the response body containing two JSON objects:

```

{"type": "profile", "verb": "post", "id": "a20ff3e5-d51e-49bc-8198-7332..."} // Green arrow, client to server
{"type": "profile", "verb": "post", "content": {"username": "Logan", "...}} // Red arrow, server to client

```

Below the messages, the status bar shows 1 requests, 0 B transferred, and 0 |.

You can see that we have an original message from the client to the server (green arrow) that establishes a new user profile. Then, a second message from the server to the client (red arrow) that lets us know that profile was successfully created. Now let's create a room and observe what happens:

Screenshot of the Network tab in the Chrome DevTools Network panel. The 'Messages' tab is selected. A single request to 'localhost' is shown, with the response body containing six JSON objects:

```

{"type": "profile", "verb": "post", "id": "a20ff3e5-d51e-49bc-8198-7332..."} // Green arrow, client to server
{"type": "profile", "verb": "post", "content": {"username": "Logan", "...}} // Red arrow, server to client
{"type": "room", "verb": "post", "id": "310bfb5e-f159-406a-98e1-de9f1..."} // Green arrow, client to server
{"type": "ACK", "verb": "post", "content": "gAAAAAABlOv4tRTusAks4..."} // Red arrow, server to client
{"type": "room", "verb": "post", "content": "gAAAAAABlOv4t3RTB1pm..."} // Green arrow, client to server
{"type": "ACK", "verb": "post", "content": "gAAAAAABlOv4uJ7ZKPJhJtn8..."} // Red arrow, server to client

```

Below the messages, the status bar shows 1 requests, 0 B transferred, and 0 |.

That looks odd, what is that id field for the "type": room message from the client? Let's copy that and try entering it on another client. The id is: '310bfb5e-f159-406a-98e1-de9f1ef70c00'. We try entering that into the other client but it does not work. Let's copy the actual room id from the first client: f5ecf61b-5bf1-468d-aaa4-7b2ba5d1d499. Obviously the code we copied was not actually the roomid. Additionally, no other message looks suspect for containing the roomid. This is

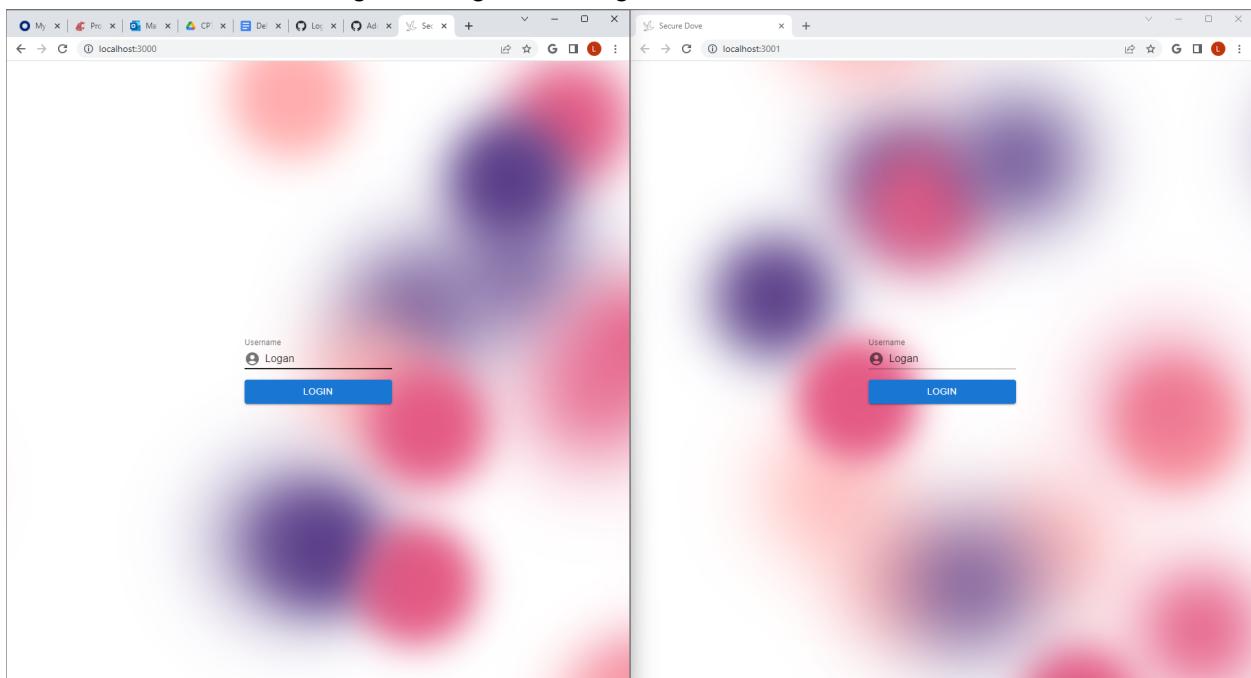
because our application encrypts the room id so that if it's intercepted, it can't be reasonably used by an attacker.

Integrity

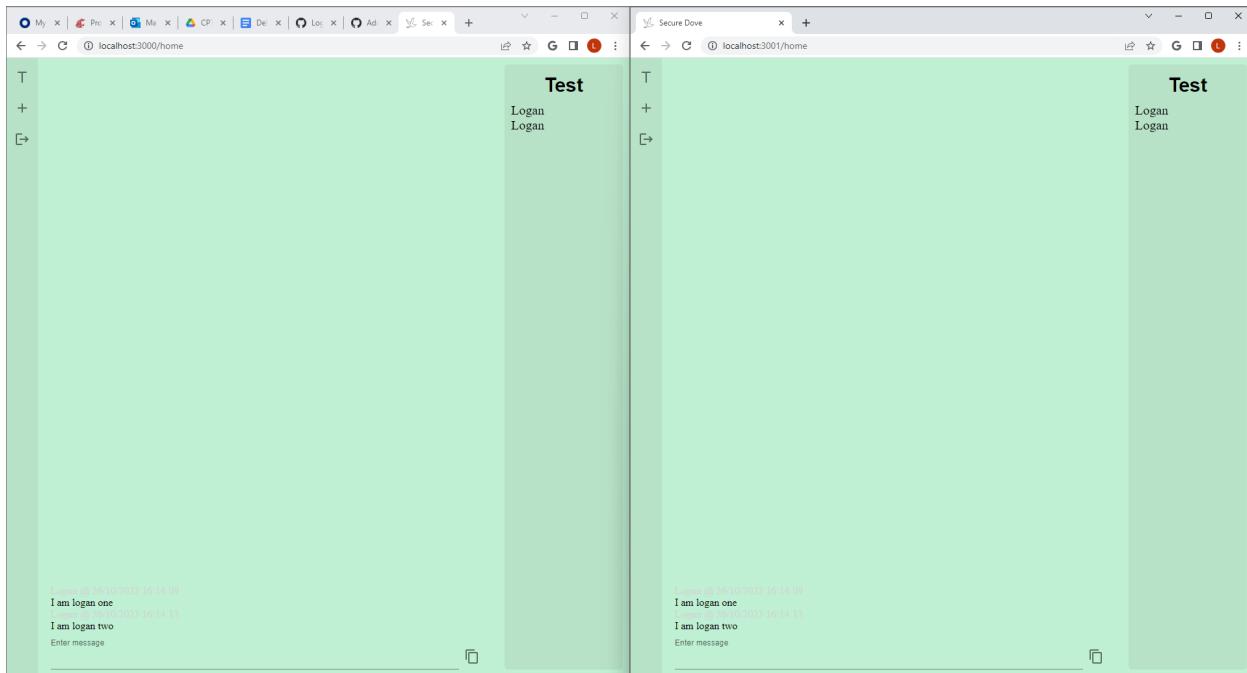
Goal: Prevent attacker from mimicking somebody else X

How: Attempt to log in with the same username on two separate clients. Join the same room and send messages.

We have two clients running, both sign in as Logan.



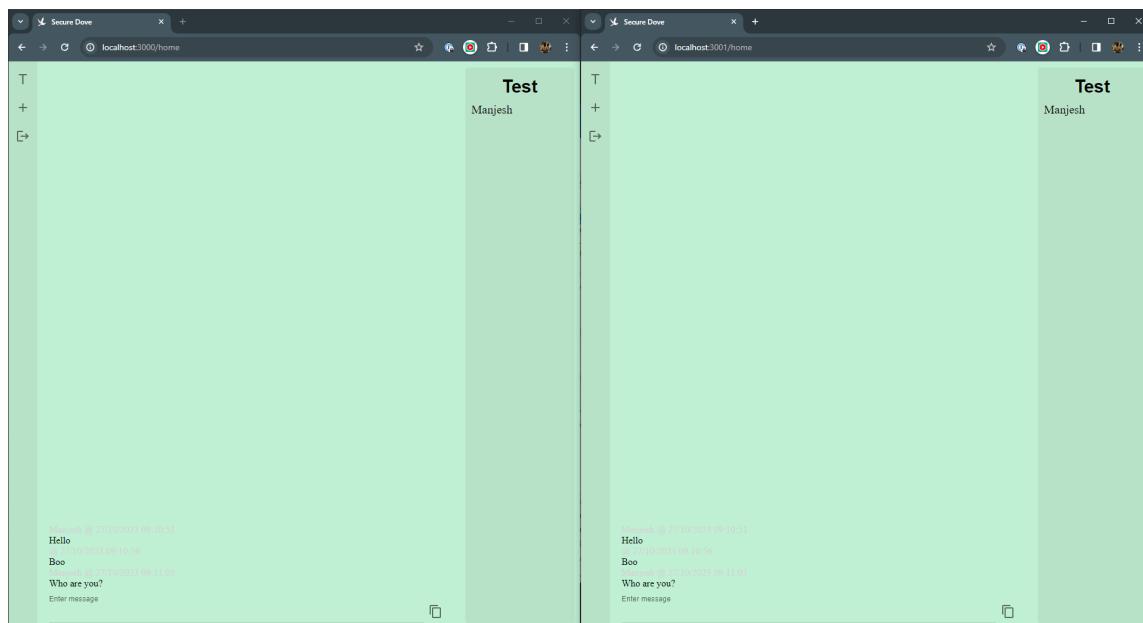
The left client creates a room which the right client joins. On the left client we also join the room, and then continue to enter the message "I am logan one". Similarly, on the right client we enter, "I am logan two".



You will notice that the message uses the same name and differs only in timestamp. Also, the user list contains two identical usernames. We can obviously see that there are two users with the same name, but to a third party we would be unable to tell which Logan had sent which message. This test proves that we need to protect against impersonation.

Goal: Transparency of who is able to see messages in a chat room X

How: Create a user with a blank name, join a chat room, can users see a visual indication that another user has joined the room?

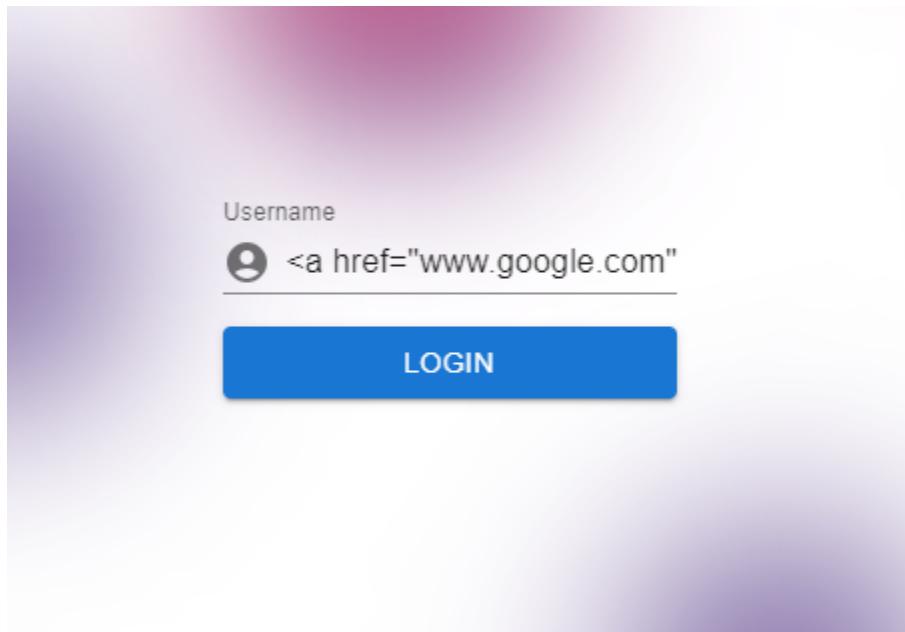


Here in this image we can see that there are two users. The one on the right is signed in with the user "Manjesh" and the one on the left is the user signed in with the user " ". This poses an issue because in the list of users, we aren't able to clearly see that another user has joined the room. Another issue is that since there is no way to tell if a user with a blank name has joined, is that you could be typing in passwords or other secure messages all while the blank user is observing everything that unfolds in front of them.

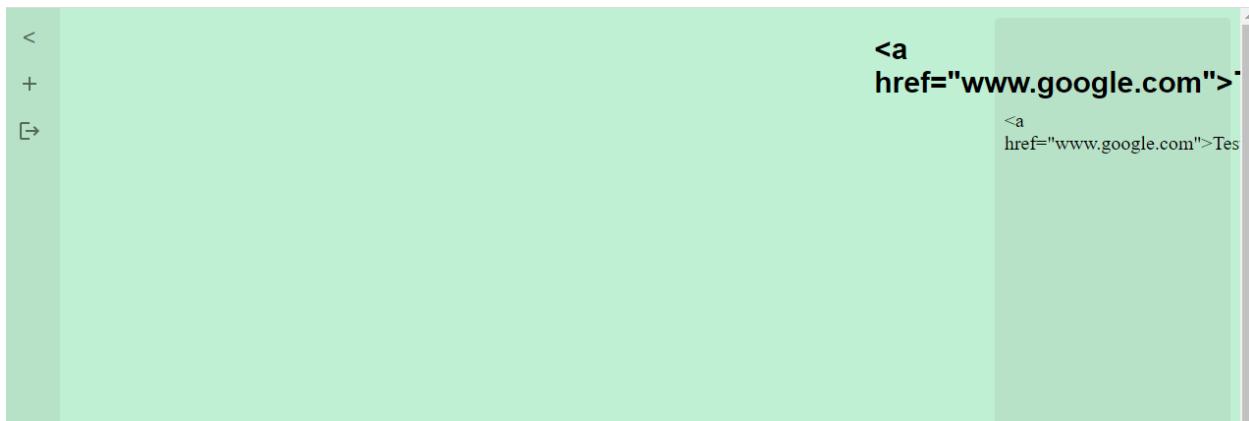
Goal: Input is not vulnerable to injected html 

How: In input elements, enter valid html to see whether a change is caused

We enter our username as an anchor tag that links to www.google.com



We create a room using the same anchor tag

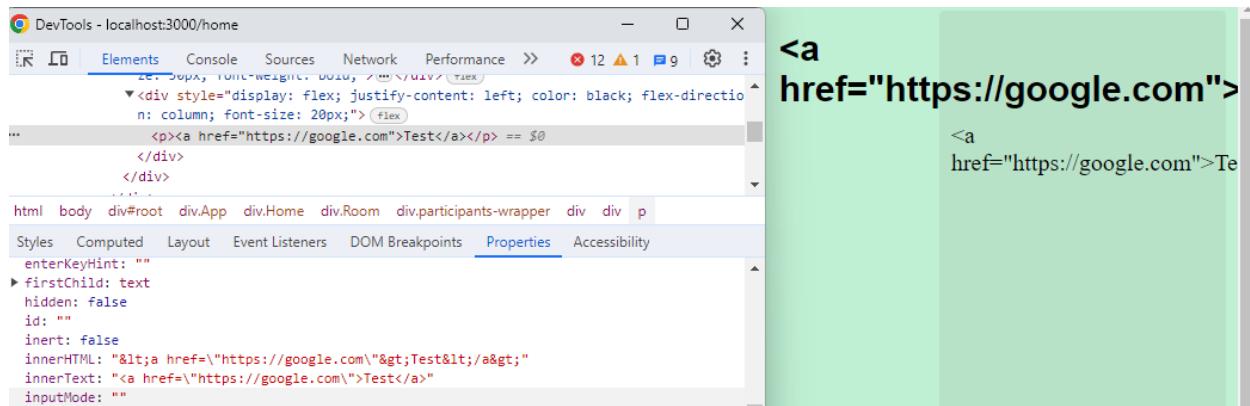


We enter a message using the same anchor tag

```
<a href="www.google.com">Test</a> @ 26/10/2023 16:09:19  
<a href="www.google.com">Test</a>
```

Enter message

We view the user as an example to see how the content is being rendered



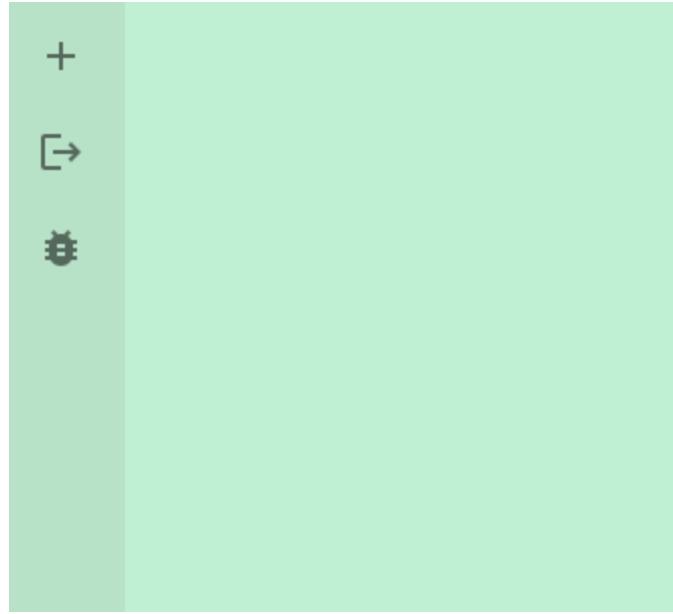
Look at the innerHTML property. You can see that special characters are correctly escaped. In this case our application correctly functions and passes this minimal assessment for vulnerability to XSS.

Availability

Goal: Malformed data in websocket does not crash server ✓

How: Send a request that the server is not able to decrypt

To test this we use a test page that can be accessed by clicking the bug icon:



Test Page:

<https://github.com/LoganKloft/LMMT-SecureDove/blob/main/client/src/components/TestClient.jsx>

Test sending a message that is unencrypted

Test sending a message that is encrypted but not in expected format

Go back to the home page

Then we send a message that is not encrypted using the first ‘send’ button. The server expects that messages are encrypted once a user is logged in. The image below shows the server output when we click ‘send’. Notice that the server detects an error and logs ‘content’. In this case, the server continues to function even if the content it receives is not encrypted.

```
IN PARSER
{'none': 'none'}
ERROR
token must be bytes or str
[]
```

How: Send a request that the server does not understand (but is able to decrypt)

For this method, we click the second ‘send’ button and observe the server behavior.

```

IN PARSER
gAAAAAB1O_XFANQ-wwPz3h-Epoj1iG9Um1UREu-Rta3CmwUbuYc06FleaWLGch_gPVqfhRYBz417wqCDtWc1ucOKm1scQelod6mhIuVM4Ez2H77VjC11M6GfnL068gSEJGrUZIEUoC0x
FINISH DECODE
{'type': 'room', 'verb': 'post', 'id': 'c4665dcf-2b4d-4e64-ba79-11e90039009e', 'content': {'HAHA': 'I AM NOT WHAT YOU EXPECTED'}}
===== SEND ACK ENCRYPTED =====
{'type': 'ACK', 'verb': 'post', 'content': {'id': 'c4665dcf-2b4d-4e64-ba79-11e90039009e'}}

===== REQUEST =====
who: logan
{'type': 'room', 'verb': 'post', 'id': 'c4665dcf-2b4d-4e64-ba79-11e90039009e', 'content': {'HAHA': 'I AM NOT WHAT YOU EXPECTED'}}

ERROR
'roomname'
[]
```

We can see that just like last time, the server catches the error and logs a message. However, the server does not crash and continues to listen for incoming requests. In fact, you can click the send buttons multiple times and then click the ‘home’ button and continue on to create your own chat rooms. This demonstrates that the server continues to work even under unexpected circumstances.

How: Send a very long request (can encryption handle long messages appropriately)

For this test we create a room called “test”. We have two clients in the room “Logan” and “Matthew”. “Logan” will send a message that contains 1000 characters which is considered a long message. We will observe the behavior on the server and also whether the message is broadcasted. If the server goes down, this is an availability vulnerability, if the message is not transmitted, it doesn’t indicate an availability vulnerability because the vulnerability only affects the message sender and is triggered by the message sender.

Retreat

Logan
Matthew



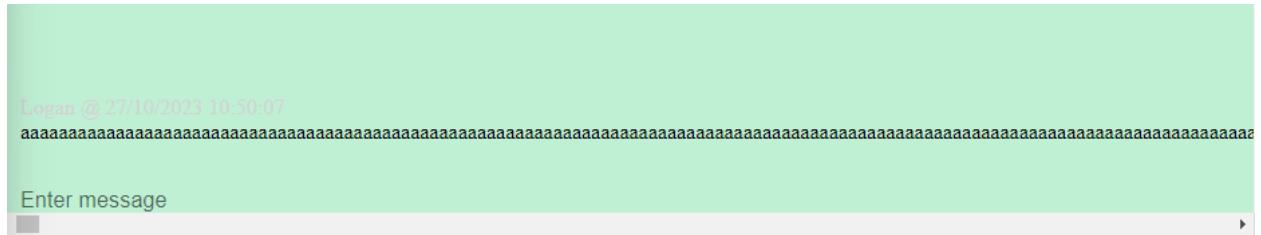
The message is displayed on the client. Let's check the server.

The server was able to parse and transmit the message to “Logan” and “Matthew”. Let’s test 10,000 characters.

Example of encrypted message on server:

```
pMp4iv79H-WH1yxFRpaIgyiCUY9SYHcwuoMkZuG7VHITgHS9n6c3ZvaJSxFLmtZycIwsl2Yw1caI-my00ccgvnljdjoavq9PwLfmzdFBEGBxw547tM9m1YX15_cRkwxa9_Pwg5yzDx500nge
jn1z5svb13dPSD4JJYcm4rnH04k16nmv_pRz6F8ZUehf9kT5qhoenxI_SnvncwQK2UExUIPqnHcn6maTCxFxeuzC5k_9evovah1vMord6ztiTQWGXHMP1k8gYqsJDMQtEKQjsK3FUE-26xt1zP
LSKqkB5wvHCE_NIPJpg1-3avAnMsuh18U1CbkqLewJLNjqYbu11_lwn9dnrgV2614D7wjhqEp3EptScuweUD01D_729ge_61NpiAvu1Xu9104c_Uwfyeorza11fd50MphvEEdwSf1z
00ULR76Ry1VM2luLS6y3Ihv6rN6Hk-3-bmRnpnbH2xaYkuE3dkcdMxHdYZ75FtyXvuhQw9NsPm82e1Leox6f16ec_rH3v0gpf-KMmcuv5zgjCM34x1oujShnAogacRp49N1umgBmmmbIRkgSP9wQNmktJNuZITClLnCA
E04t0D60ctVGrA3GSpw0OK66kXm-Oflpuv832lB1Zwjp49wlBygk760tqIBkT9Z8vF3_rH3v0gpf-KMmcuv5zgjCM34x1oujShnAogacRp49N1umgBmmmbIRkgSP9wQNmktJNuZITClLnCA
zB7M-NtR7zBzQPUhQ_C0-5jdojArNxw_c_VT75mZgocAVrq71dNA6r47tDmlhW-SfjkFytisFa-KyxmgXUSR9797qKo2dvET9P-V0vr-ZlhY3qahc7LBjVWbd7cq0208FBk-rdvlSxnhPs
zhLsK8SY2DsvlmsFtaI66W_e0My7_84fH1jjeFVsISPVlwQuCg1uTQZEvRaOfm1TmCLMeYmcl_fejkaW2Q009apKgztU99RT1eoia_OnW7oh0pHtmMzX5_oY3ryUyB3u5df51ccQk97hKuapg4
Wg6SmwNr1svvBByxHK6GeClkWRR_6vQzS2_SIkdqIavgaPRVU2RnPcQgf6Cf0ommChv7QDQdLYqj1k8AKRng7LPbnWc32lmT41f0-haSu9knqbcnMckGmxRwxKQp4CR8TTjs253dlSNtc
X4QmPm7aYIG5u9Lc192yvmHRumrF-tGoyMcuD_sT8Fcdbh8R6pH42-TtzoignC5UMJ3v-3owBsesz247mhx9aeKwqlzj3Hi35rz2mBLKtw1MyMMI_OET3DCnQ0mR_T4WbatB5baz87R67
vjCatv1QUN3CeD_wOr7n50mgmuZgsFwQ1_fapUxuGD0Uke3BVymw0g2isZt50xU_rs984eq2qjzzD8y4Pe1aXquq9o0xuCY941tgFzqV5D0-IxCsj1y30aF18Qy135VnwlaTfValh24AwWB_gL
jv64PsVGrekTPDZph6_u2rAbf5VaTgdPw0bOpGQavdm0hzkIPSWw3pEPE3bjxsrdifgbHivYlwfaSk5pvreLhfczGDIupq92A0Rgr4MqBLqQ2FpEManYZUS1PAuZGc491xeA5i-mrgNDBK
gPnIH7qLl1RuHKUkvCPUdqJbwqlkdim_5t0gjn28e1hpZWGDq520c_6peBDmd23PVKv28J0qExAZHu243PT_fltB_-qgCWYLEM_EEPA7Eda_C8-7jw_bhANKh1W4dG11DRqrJrVP3fm_xy
U7IfsRNT19m0jxz1q1x2v_LcullQ1ewvoRaLT-e0UjY5Zx_0q87UldfhAnhpzVsk59_mmU-iV5v2yBwjsveq7-V5zxdmoeRyo13ZFDZLvcj_c_u98uNPFVqjkoqxsGVE_Bh5Lucm1Rkt_ZwBA
-8t7dwvtM0tPwon4m7B-5b60v8Cx0d3EwtR3PnP4GzdsM8vLwh0J9kr3dBpDjZagJkbc1r-G5YfqoAucFabwIGo0y6gdy9mnGghiv2Fbdx6XvxxIa2rjjysapyOeGnv12JUFsnnhxwD1py
rij8PTQvx21qR2ixjUTg6K-HAJS1cK0b-3ASKNc0-qUjzdhGoeBPxFhltD4h_Qs11v1Cv0bvaQvhYUhwByF-vIYTof6N5Gijhyd_xUKMfgj7vXWVfoQHggi37Dgt98Abyl4djvI4MDa
bzyCmoLDpbI0BQ2aFWA_5GALqDhsSlbw1ym01CIw1JMFURvKExyPwUHjIn04sYFSAJjQdy1qAmErshz_cMq2B1EyAqVsS_CL_8cnHudjA6Y1IN4SwCcotzKoHp-M120KcfUtwtxF1d1k
QLqrK4D0UmEL3-8leThstzYUrtEWT2o7zKhnDrPnt2te6x7F7Ud2djjzuy5iguP2C8zjpuScnwzq1Vaz3e6gk-BzHInEejE0qg0s6xbm_AeRvtzg_ATyhLL6D0t4irrE6y0
ShK7oPg1P_y8Kv-57C81mByhR87thrXkmCx3tew2u1AmdTp7wMGXTnk0ek0CQ1CfwmlN3q1fr1-1gYFM5VkfwrPVa1yeJJTpBijn_FdM36rUymF89hfjvQgUtnIs47Fw-WzBtE9eDtq
```

What the client sees:



In both cases, we see that the message was successfully transmitted. This test passes - the encryption algorithm is able to handle large messages.

Goal: Server can detect and respond to when a message has failed to reach the client X

Metric: ACK ratio

How: modify client to arbitrarily ACK a message from the server or not

In the normal scenario we expect the ACK ratio for server responses to be 100%. To test this, we login, create a room, and send two messages. Then we compare the responses_sent to the responses_acked for the final time stamp in the audit.txt file.

```
time: 27/10/2023 10:56:24connections: 0requests_received: 0requests_acked: 0responses_sent: 0responses_acked: 0time: 27/10/2023 10:56:29connections: 0requests_received: 0requests_acked: 0responses_sent: 0responses_acked: 0responses_sent: 0responses_acked: 0time: 27/10/2023 10:56:34connections: 0requests_received: 0requests_acked: 0responses_sent: 0responses_acked: 0time: 27/10/2023 10:56:39connections: 1requests_received: 0requests_acked: 0responses_sent: 0responses_acked: 0responses_sent: 1responses_acked: 0time: 27/10/2023 10:56:44connections: 1responses_received: 0responses_acked: 0responses_sent: 0responses_acked: 0responses_sent: 1responses_acked: 0time: 27/10/2023 10:56:49connections: 1responses_received: 2responses_acked: 2responses_sent: 3responses_acked: 2time: 27/10/2023 10:56:54connections: 1responses_received: 4responses_acked: 4responses_sent: 5responses_acked: 4time: 27/10/2023 10:57:04connections: 1responses_received: 4responses_acked: 4responses_sent: 5responses_acked: 4
```

Notice that the ratio is 80%. This is because the client does not acknowledge the first response by the server (which is the response that contains the symmetric key the client should use when talking to the server). So in reality, for the messages that are supposed to be ACKed we have 100% ACK rate. Next, we drop every other ACK that the client is supposed to send. Ideally, if

there is a resolution technique, the ACKs would be 4 like before. However, if not then they will be 2. This scenario symbolizes an attacker sitting in between the client and server selectively dropping messages.

Code for dropping messages found in:

<https://github.com/LoganKloft/LMMT-SecureDove/blob/main/client/src/components/Hallway.jsx>

```
> audit.txt
time: 27/10/2023 11:14:01connections: 0requests_received: 0requests_acked: 0responses_sent: 0responses_acked: 0time: 27/10/2023
11:14:06connections: 1requests_received: 0requests_acked: 0responses_sent: 1responses_acked: 0time: 27/10/2023 11:14:11connections:
1requests_received: 2requests_acked: 2responses_sent: 3responses_acked: 1time: 27/10/2023 11:14:16connections: 1requests_received:
4requests_acked: 4responses_sent: 5responses_acked: 2time: 27/10/2023 11:14:21connections: 1requests_received: 4requests_acked:
4responses_sent: 5responses_acked: 2
```

We can see that the responses are half the amount as before. This test fails, the server is not able to respond to missing messages.

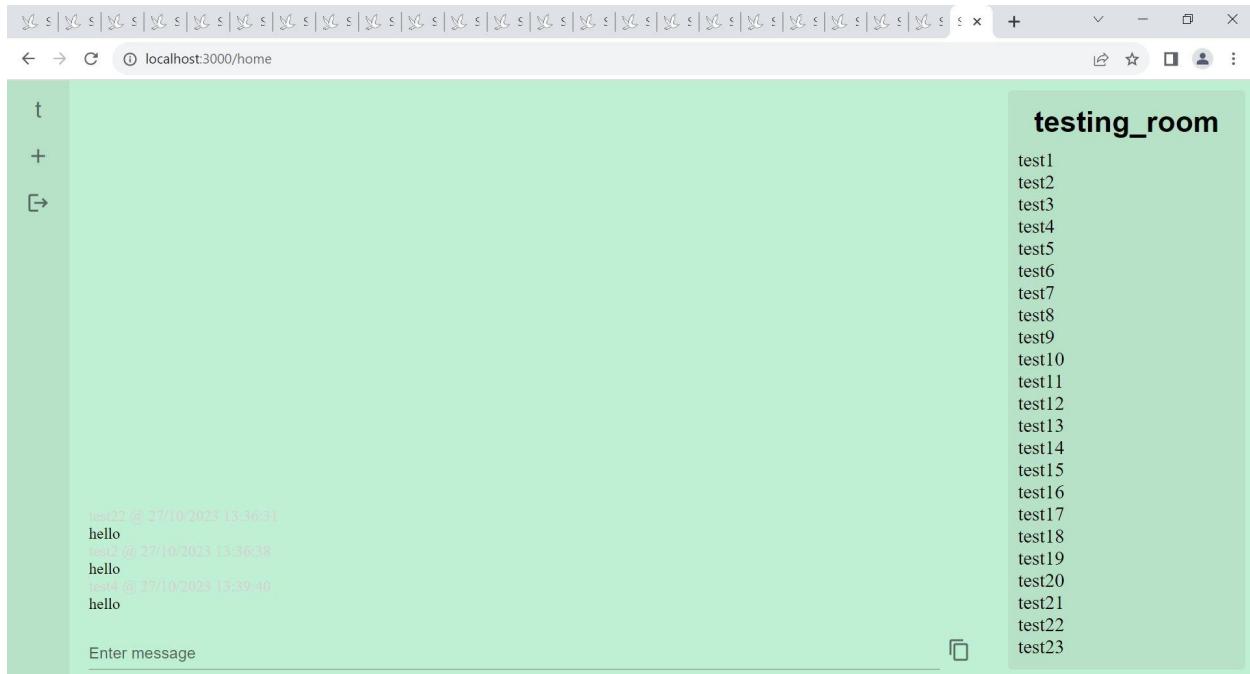
Goal: Server maintains stability when receiving too many websocket requests

How: connect a bunch of websockets to the server

To test the server for stability it needs to have a lot of clients connected or connecting, so for this test we connected 22 clients to start and placed them all in a room and sent a few messages to ensure everything was working.



This is a good amount of sockets for the server to handle so we next added client “test23” to see if the server would accept this client. It did and we added this client to the room to see if the server would also place the client in the room, which it did.



When a program was written to automate the process an error of “timed out during handshake” occurred meaning the process would have to be done manually. Which as stated above still wouldn’t have errors handling the incoming traffic.

Improved Security Plan

During the break process we observed that many of our goals withstood the tests. However, we also saw that not all security goals were met. There were three main issues. The first is that there is a window of opportunity during the fernet symmetric key transfer where the key is public. The second is that there are no defenses against users abusing the name system. A user can have the same name as another user or even have no name at all, making them silent observers in a room. The final issue is the lack of robustness for retransmitting messages that were never acknowledged by either the server or client.

Due to the present vulnerabilities in our application, we need to discuss new security plan goals and how we will address them. The following details three goals that are either new, or are an improvement to a previous goal.

Goal: Key sharing is performed such that all keys involved are not publicly discoverable. Note that this only applies to the main branch that uses fernet encryption. The asymmetric branch, which encrypts chat messages, does not share this problem.

Solution: Use private key encryption to share the symmetric key. This can be overlaid over the current process. The asymmetric branch contains a solution for key sharing that can be reused on the main branch.

Goal: Users are not able to impersonate another user or covertly observe a chatroom while inside of the chatroom.

Solution: Usernames will be limited to a minimum and maximum size. Usernames will be limited to the available character set. Usernames will be unique. These requirements must be verified on the server at the least, and preferably also on the client side for instant feedback. Creating a new user must be an atomic operation on the PROFILES object to prevent race conditions.

Goal: Messages that are ‘dropped’ will have a resolution mechanism for attempting to ensure messages are reliably transmitted.

Solution: The client and server will store a list of messages they send. Each message will contain a unique identifier. When the client or server receives a message, they will send an ACK for that message in addition to storing the message id in a container of received messages. When the client or server receives an ACK they will remove the stored message from the message container. At a particular interval, the client and server will check messages that they have not received ACKs for and rebroadcast the message. If a client or server receives a duplicate message (contains the same id) they will ACK the message but otherwise not act on it (since it is assumed that the client or server had previously acted on it). There is an opportunity to configure how many times a message should be re-broadcasted before discarding it.